

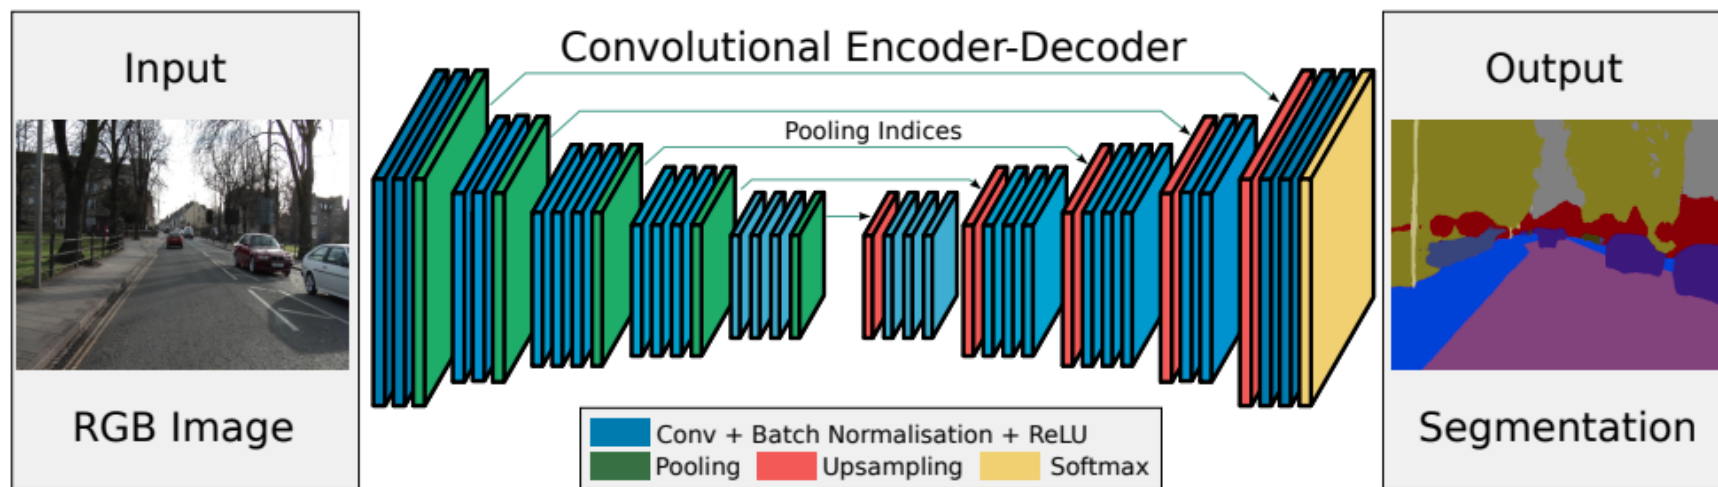
Introduction to Semantic Segmentation

<https://arxiv.org/pdf/1511.00561.pdf>




SegNet: A Deep Convolutional Encoder-Decoder Architecture for Image Segmentation

Vijay Badrinarayanan, Alex Kendall, Roberto Cipolla, *Senior Member, IEEE*,



Qualcomm sample video:
<https://www.youtube.com/watch?v=hGrJ3zuuvRQ>



DeepLab for Dense Pixel Labeling Semantic Image Segmentation

<https://github.com/google-research/deeplab2> and older one <https://github.com/tensorflow/models/tree/master/research/deeplab>


DeepLab2: (1) a TensorFlow library for deep labeling for a unified and state-of-the-art TensorFlow codebase for dense pixel labeling, including, but not limited to semantic segmentation, instance segmentation, panoptic segmentation, depth estimation, or even video panoptic segmentation. (2) Deep labeling assigns a predicted value for each pixel.

Waymo Open Dataset: Panoramic Video Panoptic Segmentation

Waymo Open Dataset: Panoramic Video Panoptic Segmentation

Jieru Mei^{1*} Alex Zihao Zhu² Xincheng Yan² Hang Yan²
Siyuan Qiao³ Yukun Zhu³ Liang-Chieh Chen³
Henrik Kretzschmar² Dragomir Anguelov²



¹Johns Hopkins University ²Waymo LLC ³Google Research



[cs.CV] 15 Jun 2022

Semantic Decoder Python

1. Pointwise convolution, e.g., 1x1xk convolution;
2. Depthwise convolution: Atrous convolution;





<https://www.analyticsvidhya.com/blog/2019/02/tutorial-semantic-segmentation-google-deeplab/>

Atrous convolutions

$$y[i] = \sum_k x[i + r \cdot k] w[k]$$


DeepLab uses atrous convolution with rates 6, 12 and 18.



Dilated Convolution

Standard Convolution

Convolution Up-sampling



<https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>

1

Nearest Neighbor

1	2
3	4

Input: 2 x 2



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Output: 4 x 4

3

Bi-Linear Interpolation

10	20
30	40

2x2

2x

10	12	17	20
15	17	22	25
25	27	32	35
30	32	37	40

4x4

3

"Bed of Nails"

1	2
3	4

Input: 2 x 2



1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

Output: 4 x 4

<https://towardsdatascience.com/transposed-convolution-demystified-84ca81b4baba>

4

Max Pooling

Remember which element was max!

1	2	6	3
3	5	2	1
1	2	2	1
7	3	4	8

Input: 4 x 4

5	6
7	8

Output: 2 x 2

... Rest of the network

Max Unpooling
Use positions from pooling layer

1	2
3	4

Input: 2 x 2



0	0	2	0
0	1	0	0
0	0	0	0
3	0	0	4

Output: 4 x 4


Transposed Convolution Up-sampling

Credit of the example illustration:
<https://towardsdatascience.com/transposed-convolution-demystified-84ca81b4baba>


<https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>

1. Consider a 2x2 encoded feature map which needs to be upsampled to a 3x3 feature map.

Input image Feature map: 2x2
 Kernel 2x2
 Upsampled output image: 3x3




kernel of size 2x2 with unit stride and zero padding





Step 1. Feature map and the kernel

Step 2. Transposed convolution for each pixel in the feature map



Step 3. Add output at each pixel location together to form upscaled image



Transposed Convolution Up-sampling with Python

Python Reference Code (Untested)


```
def apply_conv(data, kernel, conv):
    """
    Args:
    data (NDArray): input data.
    kernel (NDArray): convolution's kernel parameters.
    conv (Block): convolutional layer.
    Returns:
    NDArray: output data (after applying convolution).
    """
    # add dimensions for batch and channels if necessary
    while data.ndim < len(conv.weight.shape):
        data = data.expand_dims(0)
    # add dimensions for channels and in_channels if necessary
    while kernel.ndim < len(conv.weight.shape):
        kernel = kernel.expand_dims(0)
    # check if transpose convolution
    if type(conv).__name__.endswith("Transpose"):
        in_channel_idx = 0
    else:
        in_channel_idx = 1
    # initialize and set weight
    conv._in_channels = kernel.shape[in_channel_idx]
    conv.initialize()
    conv.weight.set_data(kernel)
    return conv(data)
```

<https://medium.com/apache-mxnet/transposed-convolutions-explained-with-ms-excel-52d13030c7e8>

Unet For Semantic Segmentation

<https://towardsdatascience.com/understanding-semantic-segmentation-with-unet-6be4f42d4b47>


Unet For Semantic Segmentation



Code sample and tutorial:
<https://pyimagesearch.com/2022/02/21/u-net-image-segmentation-in-keras/>

Unet For Semantic Segmentation

<https://towardsdatascience.com/understanding-semantic-segmentation-with-unet-6be4f42d4b47>



GIPHY and Other Tools for Annotation of Images for Semantic Segmentation

<https://giphy.com/gifs/R0dnXaKJowlR2yL5CG>

[Labelbox](#)

[Supervisely](#)

[Fritz AI](#)

[RectLabel](#)

[Anolytics](#)

[Playment](#)

[Appen](#)

[Scale.ai](#)



<https://cnvrg.io/semantic-segmentation/>

Six (6) Useful Image Segmentation Datasets And Python Usage

<https://cnvrg.io/semantic-segmentation/>

coco:

```
import tensorflow_datasets as tfds
(X_train, X_test), ds_info = tfds.load(
    'coco',
    split=['train', 'test'],
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)
```

waymo:

```
import tensorflow_datasets as tfds
(X_train, X_test), ds_info = tfds.load(
    '~waymo_open_dataset',
    split=['train', 'test'],
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)
```

PASCAL:

```
import tensorflow_datasets as tfds
(X_train, X_test), ds_info = tfds.load(
    '~voc',
    split=['train', 'test'],
    shuffle_files=True,
    as_supervised=True,
    with_info=True,
)
```